

**Politecnico di Milano**  
**Facoltà di Ingegneria dell'Informazione**



**Corso di Laurea Magistrale in Ingegneria Informatica**  
**Dipartimento di Elettronica, Informazione e**  
**Bioingegneria**

**... Titolo della tesi ...**  
**... al massimo su due righe ...**

**Advisor: Emanuele DELLA VALLE**  
**Co-Advisor: Daniele DELL'AGLIO**

**Master thesis by: Riccardo TOMMASINI matr. 799120**

**Academic Year 2013-2014**



*dedica...*



# Thanks To

Ringraziamenti vari, massimo una o due pagine.

Milano, 1 Aprile 2005

... *Fabio* ...



# Estratto

abstract in italiano





# Abstract

abstract in english



# Table of Contents

<b>1</b>	<b>Introduction</b>	<b>1</b>
1.1	Contributions . . . . .	1
1.1.1	Heaven . . . . .	1
1.1.2	RSPEngine inside DSMS . . . . .	1
1.1.3	Experiments . . . . .	1
1.2	Structure of this Thesis . . . . .	1
<b>2</b>	<b>Background</b>	<b>3</b>
2.1	Semantic Web . . . . .	3
2.2	Stream Processing . . . . .	3
2.3	Stream Reasoning . . . . .	3
2.4	Empirical Research . . . . .	3
2.4.1	Software Testing . . . . .	4
2.5	Benchmarks . . . . .	4
2.5.1	TCP . . . . .	4
2.5.2	Reasoning Benchmark . . . . .	4
2.5.3	DBMS & CEP Benchmarking . . . . .	4
2.5.4	Stream Reasoning Benchmark . . . . .	4
<b>3</b>	<b>Problem Settings</b>	<b>7</b>
3.1	Comparative Research . . . . .	7
3.2	Requirements . . . . .	10
<b>4</b>	<b>Conclusions and Future Works</b>	<b>13</b>
<b>5</b>	<b>Conclusions</b>	<b>15</b>
5.1	Research Question . . . . .	15
5.2	Results . . . . .	15

## TABLE OF CONTENTS

---

5.3	Limitations . . . . .	15
5.4	Future Works . . . . .	15
<b>Appendices</b>		<b>17</b>
<b>A</b>	<b>...</b>	<b>19</b>
A.1	Introduction . . . . .	19
A.2	... . . . .	19
<b>B</b>	<b>...</b>	<b>21</b>
B.1	Introduction . . . . .	21
B.2	... . . . .	21
<b>Bibliography</b>		<b>23</b>

# Chapter 1

## Introduction

### 1.1 Contributions

#### 1.1.1 Heaven

#### 1.1.2 RSPEngine inside DSMS

#### 1.1.3 Experiments

### 1.2 Structure of this Thesis

Blablabla

Esempio di figura

Figure 1.1: Didascalia esempio di figura.

Esempio di equazione

$$\dot{V}_4 = -k_v \lambda (z_1 - \varepsilon \lambda z_3)^T (z_1 - \varepsilon \lambda z_3). \quad (1.1)$$



# Chapter 2

## Background

### 2.1 Semantic Web

### 2.2 Stream Processing

### 2.3 Stream Reasoning

### 2.4 Empirical Research

Tichy and collaborators [15] evaluated 400 articles published in 1993, 50 of them randomly selected papers published by ACM in 1993 and the rest systematically selected from a few journals in Systems and Software Engineering, and classified the research reported in the paper in five categories (quoting [15] definitions):

- Formal theory: articles whose main contributions are formally tractable propositions, e.g., lemmata and theorems and their proofs.
- Design and modelling: systems, techniques, or models, whose claimed properties cannot be proven formally. Examples include software tools, performance prediction models, and complex hardware and software systems of all kinds. The papers in this class were further classified in the categories 0%, 0-10%, 10-20%, 20-50 and +50%, according to the proportion of the paper that was dedicated to the evaluation of the new system, technique, or model.

## Background

---

- Empirical work: articles that collect, analyse, and interpret observations about known designs, systems, or models, or about abstract theories or subjects (as this paper does). The emphasis is on evaluation, not on new designs or models.
- Hypothesis testing: articles that define hypotheses and describe experiments to test them.
- Others: articles that do not fit any of the four categories above, e.g., surveys.

### 2.4.1 Software Testing

SOAK

Stress

## 2.5 Benchmarks

### 2.5.1 TCP

### 2.5.2 Reasoning Benchmark

LUBM

### 2.5.3 DBMS & CEP Benchmarking

Linear Road

### 2.5.4 Stream Reasoning Benchmark

SRBench

SRBench [5] proposes a suite of test queries and defines metrics to evaluate the performance of the system. This benchmark contains 17 queries to gather the properties of the RDF stream engines. The queries vary to ensure that several features of the target system are tested: queries involving single or multiple input streams, queries over stream-only data sources or over mixed stream and static data source, etc. In [5] the authors applied the benchmark on the existent RDF stream engines, and explained the differences in term of



supported functionalities. Time and memory performance tests, and scalability tests are not targeted in the actual version of SRBench.

### **LSBench**

LSBench [6] proposes three tests to evaluate the RDF stream engines. The first one is a functional test to verify the operators and the functionalities supported by the engines: it is a test similar to the one proposed by SRBench. The second test is a correctness test: its goal is to verify if the tested RDF stream engine produces the correct output. Actually this analyses only the number of produced answers, assuming that the contents of the output are correct. Finally, the third test is a maximum input throughput test: it has the goal evaluate the maximum throughput of the RDF stream engines. This test is done increasing the rate of data in the stream and verifying the number of the answers. For each test a set of 12 queries is provided; similarly to SRBench, the queries vary to take into account different features of the engines (single and multiple streams, presence of static data, etc)

### **Correttezza**

### **Seven Commandaments**



# Chapter 3

## Problem Settings

In this chapter we state the questions that leads our research and the issues it poses. We present a possible solution and the requirements it must satisfy.

### 3.1 Comparative Research

The Computer Science (CS) community classifies its activity in many works and shows that its research mostly follow an engineering epistemology [?, ?]. A majority of publications belongs to *empirical work* and *design & modelling* classes of Tichy's taxonomy (See 2.5.1 ) and, among them, proposals of new systems or models are more common then the evaluations of existing ones. This contrasts with other engineering areas, where the experimental research is almost dominant. CS needs to focus on collecting, analysing, and interpreting the observations of those works it usually only designs or implements. Now we demand, *What are the motivation under this CS research lack?* The main problem in evaluating software systems or models regards their complex and multifaceted nature. Other explanations concern the difficulties of conducting realistic evaluations, because of the number of the involved variables.

Social Sciences typically use a Systematic Comparative Research Approach (SCRA), because the complexity of its research subjects goes beyond the possible observable models. For this reason, they exploit techniques to deal with complex cases that can not be simplified in experimental setting. The analysis of a single case study allows the deep comprehension, but it makes difficult to engage any form of generalisation. On the other hand, cross-case studies are more relevant and allow general thinking, but their final complexity represents

## Problem Settings

---

a problem. We need a strategy that reduces the analysis complexity without lose the relevance of each involved system. In this regard, Russel Schutt discusses four stages to systematic qualitative comparative studies for history social phenomena:

S.1 Premise of the investigation: identification of possible causes.

S.2 Choose the cases to examine (location, language, gender).

S.3 Examine the similarities and the differences with shared methods.

S.4 Propose a causal explanation for the phenomena.

Moreover, complex cases are seen as a combination of known properties upon which is possible to identify parallelism or state contrasts and are used to set up experimental environment. We can define an *experiment* as a test under controlled conditions that is made to demonstrate a known truth or examine the validity of an hypothesis (W.R Inge). Thanks to experiments, we can exploit the notions of *reproducibility*, appreciating variations on changing experiment conditions, *repeatability*, consolidating observations trough multiple identical executions, and *comparability*, contrasting the results to identify the differences.

Case-driven analysis and the notion of experiment are not totally new concepts in Computer Science. Database community firstly explores the idea of comparative research trough benchmarking techniques and actually the quality of empirical studies is rising. It is worth to note Jim Gray’s work about transactional benchmarking (TCP, Section 2) and Domain Specific Benchmarks (DSB). He states that *any comparison on performances starts with the definition of a benchmark or a workload*, but we need know the relevance of the metrics. Measurement variations are very frequent from one application to another, because each system is thought to solve a small problems set. A DSB must response properly to system diversities, by specifying a synthetic workload to describe typical applications in the problem domain; moreover it must provide workload performances on various systems and an estimation of relative performance on the problem domain. Gray proposes also four criteria that a DSB must meet, which are:

G.1 RELEVANCE, it must measure the performance peak of systems when performing domain typical operations.

G.2 PORTABILITY, it must be easy-to-implement on many different systems and architectures

G.3 SCALABILITY, it must be meaningful for both small and large computer systems

G.4 SIMPLICITY, it must be understandable to obtain credibility

We note a relation between Schutt work in social science and Gray’s criteria: Gray, in G.1, identifies the relevant metrics for the evaluation and Schutt, in S.1, demands a pre-analysis phase of the phenomenon. Problem scaling is not so relevant in social science as for benchmarking. Case defining properties also include the problem dimension (S.2), in DB context is required by both G.3 and G.2, which demands implementation-related conditions. Last but not least, G.4 demands simplicity to obtain credibility while S.3 suggests to exploit those evaluation methods that are commonly accepted (shared) by the research community.

The growing number of RDF Stream Processing techniques alarms the Stream Reasoning (SR) community, which identifies the need of evaluation practices to allow comparison. The systems to study, RSP Engines (See ??), has an high resulting complexity: the execution mechanism they implement summed to their internal semantic and the environment where this execution happens are evaluation-related characteristics, which increase this complexity and demonstrate that a meaningful comparison between RSP Engines is non-trivial. Chapter 6 shows in details that may be difficult to analyse those system, even in case of simple and well defined architectures. The interest on cros-case analysis between complex subjects draw the need of a comparative research approach. Initial efforts in this direction try to define frameworks that resume DSMSs [??] and Reasoning [??] benchmarking studies. LS Bench and SR Bench propose a set of queries, data sets, and methods to test and compare existing RSP engines [??]. Both these works share a common background: the Linear Road Benchmark (LRB) is the only existing benchmark for relational data stream processing engines. Actually it only states which requirements a solution must satisfy, without proposing a concrete one. LS Bench and SR Bench implement and extend this work, but they do not re-contextualise LRB requirements in SR context. Following discussions identify other lacks, i.e none

## Problem Settings

---

of them completely face the problem of evaluating query result correctness, because they do not analyse engine semantics. LS Bench concentrates attention to the evaluation of engines throughput and it checks the correctness of the query result measuring the mismatch after comparing different RSP Engines. SR Bench entirely ignores this issue and analyses the coverage of SPARQL constructs for each commercial engine. More recent works describe deeply all the RSP Engine properties, identify the future challenges and provide a standardization of the meaningful stress test for RSP Engine [?]. They define benchmarking requirements, commandments on how to test RSP Engines.

In the end, LRB provides also a simulator to validate the benchmarked DSMS systems. Stream Reasoning community still lacks an infrastructure that can control the execution environment and allow to systematic testing. Previously the community questioned itself *How to support SRCA on RSP Engines?* Now we have queries, dataset and methods, that partially answer it and the new research question is ”*Does SR research needs a validator to enable the SCRA?*”. From the aerospace engineering we borrow the idea of an *Engine Test Stand*. This solution allows experiments design, their systematic execution and automatic results comparison, three ingredients to properly answer the research question.

## 3.2 Requirements

In order to simplify and support Systematic Comparative Research Approach on RSP engines trough an Engine Test Stand we need to answer the following questions:

Q.1 How can the behaviour of system be evaluated?

Q.2 What makes this evaluation rigorous?

Q.3 How can this rigorous evaluation be automated?

To answer Question Q.1 we exploit traditional definition of *experiment* presented above. We answer Q.2 referring to *reproducibility*, *repeatability*, and *comparability* of experiments. Trough this concepts it is easy to answer Q.3 formalising the requirements for the solution.

*Reproducibility* refers to measurement variations on a subject under changing conditions. We gather this conditions into experiment configuration, whose specification is up to the user. For this reason the solution must be independent from:

- R.1 *Test data*, relevant RDF data streams and ontologies chosen from user domain of interest.
- R.2 *Query*, relevant queries registered from user domain of interest.
- R.3 *Engine*, any RSP Engine tested by the means of easy-to-implement software interfaces.

*Repeatability* refers to variations on repeated measurements on a subject under identical conditions. The solution must not affect the RSP engine evaluation, which, from a practical point of view, poses two requirements:

- R.4 it must not be running when the RSP engine is under execution.
- R.5 it must have reduced (and possibly constant) memory footprint.

*Comparability* refers to performance measurements nature and the relations between experimental conditions. The SCRA demands both the definition of *comparable metrics* and the standardization of *evaluation methods*, which means the solution must:

- R.6 include *basic set of performance measurements* [1].
- R.7 enable users extensions with new software sensors and specific measurements collection.
- R.8 support performance measurements collection for further analysis.
- R.9 allow *qualitative analysis* through tools for result visualization

Previous works about Stream reasoning [] shows that the minimal performance measure set includes **Latency** – defined as the delay between the injection of an event in the RSP engine and its response to it –, **Memory Load** – defined as the difference between total system memory and the free one –, and **Completeness & Soundness** of query-answering results.

In terms of software engineering, any solution which satisfies the requirements above demands also some technical ones:

## Problem Settings

---

- R.10 *Extendible Design*, the possibility to replace theoretically each module with one with the same interfaces, but different behaviour, without affecting architecture stability.
- R.11 *Event-base architecture* to properly communicate with RSP Engines, which normally exploit it.
- R.12 *Easy-to-Parse RDF Serialisation* for the events presented to the RSP Engine in exam

SCRA is case-oriented, it needs *successful analysis and evaluations examples* to pose experimental guidelines and *initial terms of comparison*. We call baseline an elementary solution for the RSP problem, which is relevant from an experimental viewpoint. To fulfil this request the answer to our research question must consider to have specific baseline modules and to satisfy their own requirements, which are:

- R.13 It must be Elementary: requiring the minimum design effort, which means be a non-innovative solution and do not care about performances.
- R.14 It must be Eligible: being a fair term of comparison w.r.t. commercial solutions.
- R.15 It must be Relevant: covering one of the theoretical solutions.
- R.16 It must be Simple: allowing to identify easily those characteristics which support hypothesis formulation and comparison.



## Chapter 4

### Conclusions and Future Works



# Chapter 5

## Conclusions

### 5.1 Research Question

### 5.2 Results

### 5.3 Limitations

### 5.4 Future Works



# Appendices



# Appendix A

...

## A.1 Introduction

Introduzione agli argomenti trattati nell'appendice, dalle 4 alle 10 righe.

## A.2 ...

Argomenti trattati suddivisi sezione per sezione. Alla fine del capitolo non includere alcun sommario.





# Appendix B

...

## B.1 Introduction

Introduzione agli argomenti trattati nell'appendice, dalle 4 alle 10 righe.

## B.2 ...

Argomenti trattati suddivisi sezione per sezione. Alla fine del capitolo non includere alcun sommario.



# Bibliography

- [1] Thomas Scharrenbach, Jacopo Urbani, Alessandro Margara, Emanuele Della Valle, and Abraham Bernstein. Seven commandments for benchmarking semantic flow processing systems. In *ESWC*, pages 305–319, 2013.